

Video tutorials for the NSIGHT tools

- <https://www.youtube.com/watch?v=nYSdsJE2zMs>
- <https://www.youtube.com/watch?v=DLQwldhrL1A> (start at 16:20)

In depth resources about Nsight compute and Nsight systems can be found [here](#) and [here](#). The most applicable uses of these profilers are explained below.

### **Using Nsight Compute**

To profile a program:

```
Unset
ncu -o <reportName> <executable + flags>
```

For example, to run this profiler on *cudaScan*, we could do something like this:

```
ncu -o profile ./cudaScan -m scan -i random
```

Where “profile” will be the name of the report file. You should also see the profiler spit out its progress into the terminal, which will look something like this:

```
==PROF== Profiling "upstream(int, int, int *)" - 0: 0%...50%...100% - 8 passes
==PROF== Profiling "upstream(int, int, int *)" - 1: 0%...50%...100% - 8 passes
==PROF== Profiling "upstream(int, int, int *)" - 2: 0%...50%...100% - 8 passes
==PROF== Profiling "upstream(int, int, int *)" - 3: 0%...50%...100% - 8 passes
==PROF== Profiling "upstream(int, int, int *)" - 4: 0%...50%...100% - 8 passes
==PROF== Profiling "upstream(int, int, int *)" - 5: 0%...50%...100% - 8 passes
==PROF== Profiling "setZero(int, int *)" - 6: 0%...50%...100% - 8 passes
==PROF== Profiling "downstream(int, int, int *)" - 7: 0%...50%...100% - 8 passes
```

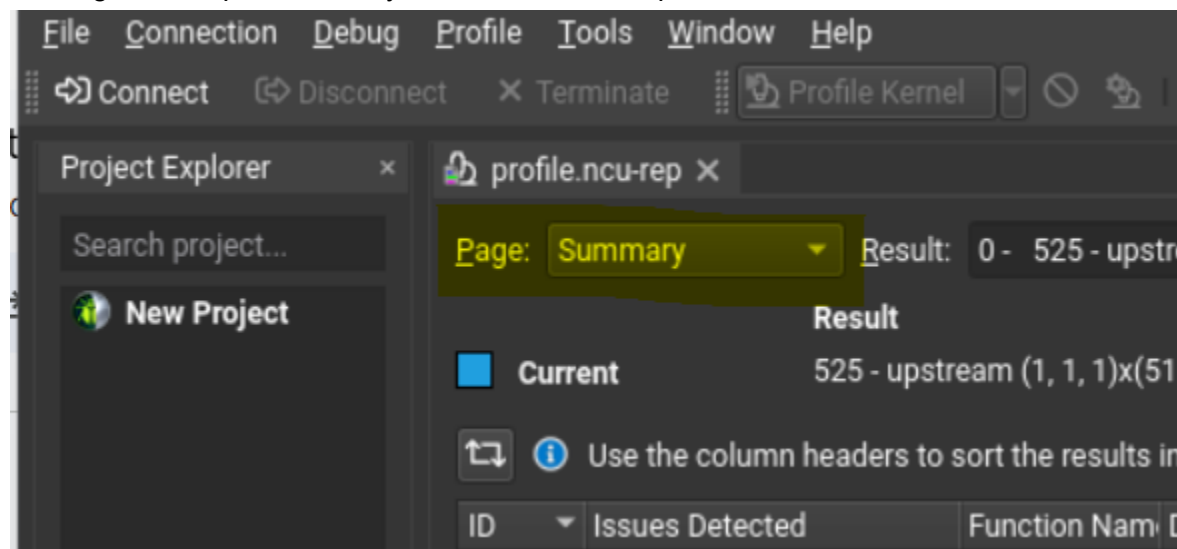
This should create a report file of the same name that you have chosen, and in your current directory you should see the file with the \*.ncu-rep file type. To view this file, we will use NVIDIA's NCU UI:

```
Unset
ncu-ui <reportName.ncu-rep>
```

Upon entering this command, a popup like this should appear:

ID	Issues Detected	Function Name	Demangled Name	Process	Device Name	Grid Size	Block Size	Cycles [cycle]	Duration
0	3	upstream	upstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,704	2,704
1	3	upstream	upstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,735	2,735
2	3	upstream	upstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,734	2,734
3	3	upstream	upstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,733	2,733
4	3	upstream	upstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,735	2,735
5	3	upstream	upstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,720	2,720
6	4	setZero	setZero(int, int *)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	1, 1, 1	2,210	2,210
7	3	downstream	downstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,736	2,736
8	3	downstream	downstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,739	2,739
9	3	downstream	downstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,740	2,740
10	3	downstream	downstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,744	2,744
11	3	downstream	downstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,743	2,743
12	3	downstream	downstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,731	2,731
13	3	upstream	upstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,711	2,711
14	3	upstream	upstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,729	2,729
15	3	upstream	upstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,722	2,722
16	3	upstream	upstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,730	2,730
17	3	upstream	upstream(int, int, int)	[9019] cudaScan	NVIDIA GeForce RTX 2080	1, 1, 1	512, 1, 1	2,717	2,717

Looking at the top left corner you should see a drop down menu like this:



Here you can choose different parts of the report that the profiler has collected. The *Details* page has the most understandable information, and definitions of terms used can be explained by hovering over the term. The *Source* page can be interesting to look at if lower-level source code is interesting to you. Feel free to explore other pages as well.

## Using Nsight Systems

Note: Conduct these profilers on terminating code, ie. use the “-b” (bench) option for the render assignment so that the program will terminate. Also note for the `./render` executable make sure to include the “-r cuda” option such that the code is run through CUDA.

To view how long specific API calls and GPU activities take, we can use another part of NVIDIA's profiling tools called Nsight systems. To get a view of this data you can use:

Unset

```
nvprof <executable>
```

The format of the executable is identical to the example used above.

The result of this command should look like this:

```
=19862= Profiling result:
      Type  Time(%)   Time      Calls      Avg      Min      Max  Name
GPU activities: 95.47% 252.18ms    65  3.8797ms  2.0885ms  31.896ms [CUDA memcpy DtoH]
                3.28%  8.6737ms    64  135.53us  134.08us  136.93us execute(void)
                1.24%  3.2770ms    64  51.202us  50.239us  52.319us kernelClearImage(float, float, float, float)
                0.00% 11.232us     9  1.2480us  1.2160us  1.4080us [CUDA memcpy HtoD]
API calls:      63.43% 266.81ms    69  3.8668ms  6.8050us  33.383ms cudaMemcpy
                31.53% 132.61ms     5  26.523ms  1.8980us  132.53ms cudaMalloc
                3.97%  16.701ms   192  86.984us  802ns    147.39us cudaDeviceSynchronize
                1.01%  4.2316ms   128  33.059us  3.9240us  105.76us cudaLaunchKernel
                0.03% 113.86us   101  1.1270us  155ns    47.808us cuDeviceGetAttribute
                0.02%  86.138us     1  86.138us  86.138us  86.138us cudaGetDeviceProperties
                0.01%  35.882us     5  7.1760us  6.9620us  7.6170us cudaMemcpyToSymbol
                0.01%  22.200us     1  22.200us  22.200us  22.200us cuDeviceGetName
                0.00%  4.8740us     1  4.8740us  4.8740us  4.8740us cudaGetDeviceCount
                0.00%  4.2790us     1  4.2790us  4.2790us  4.2790us cuDeviceGetPCIBusId
                0.00%  1.4480us     3    482ns   251ns    925ns cuDeviceGetCount
                0.00%  657ns        2    328ns   148ns    509ns cuDeviceGet
                0.00%  397ns        1    397ns   397ns    397ns cuDeviceTotalMem
                0.00%  317ns        1    317ns   317ns    317ns cuModuleGetLoadingMode
                0.00%  284ns        1    284ns   284ns    284ns cuDeviceGetUuid
```

To get a similar but slightly more in-depth view of the information, you can use:

Unset

```
nsys profile --stats=true <executable>
```

You should see something like this:

```
Overall: 0.0801 sec (note units are seconds)
Generating '/tmp/nsys-report-bee8.qdstrm'
[1/8] [=====100%] report1.nsys-rep
[2/8] [=====100%] report1.sqlite
[3/8] Executing 'nvtxsum' stats report
SKIPPED: /afs/andrew.cmu.edu/usr23/czlu/private/15418/asst2/render/report1.sqlite does not contain NV Tools Extension (NVTX) data.
[4/8] Executing 'osrtsum' stats report

Operating System Runtime API Statistics:

Time (%)  Total Time (ns)  Num Calls  Avg (ns)  Med (ns)  Min (ns)  Max (ns)  StdDev (ns)  Name
-----
62.8      330,478,984      2  165,239,492.0  165,239,492.0  1,106,422  329,372,562  232,119,213.6  sem_wait
19.0      100,059,367     13   7,696,874.4   1,647,986.0    1,513    35,333,812  11,969,494.1  poll
10.3       53,932,499     515   104,723.3    13,047.0     1,023   18,431,448   951,214.0    ioctl
7.0       36,779,015     44   835,886.7     8,250.5     1,048   19,559,912   3,840,093.0  fopen
0.3       1,557,899      31    50,254.8     4,447.0     3,052   1,091,476   194,633.9    mmap64
0.2       996,564         5   199,312.8     1,969.0     1,105   564,902    275,329.7    fcntl
0.2       853,453        10    85,345.3     65,430.5    43,271   294,651    75,055.7    sem_timedwait
0.1       398,102        49     8,124.5     7,404.0     2,132    22,229     3,913.2    open64
0.1       376,798         5   75,359.6     80,775.0    35,641   139,528    42,875.0    pthread_create
0.0       215,719        26     8,296.9     3,082.0     1,004   129,487    24,832.3    fclose
0.0       189,875         7   27,125.0     25,176.0    1,479   49,885    15,367.3    fgets
0.0       112,677        16     7,042.3     5,327.5     1,281    30,962     7,659.7    mmap
0.0       38,302          9     4,255.8     3,986.0     1,903    6,621     1,531.6    fread
0.0       32,575          5     6,515.0     5,222.0     2,017   12,571     4,036.2    open
0.0       27,435          7     3,919.3     3,489.0     1,325   10,426     3,033.9    munmap
0.0       20,458          5     4,091.6     2,234.0     1,208     9,339     3,645.7    read
0.0       12,018          1   12,018.0    12,018.0    12,018   12,018     0.0    fopen64
0.0       10,365          6     1,727.5     1,644.5     1,116    2,709     572.2    write
0.0        7,991          2     3,995.5     3,995.5     2,026    5,965    2,785.3    socket
0.0        7,938          1     7,938.0     7,938.0     7,938    7,938     0.0    connect
0.0        4,632          1     4,632.0     4,632.0     4,632    4,632     0.0    pipe2
0.0        1,266          1     1,266.0     1,266.0     1,266    1,266     0.0    bind
0.0        1,128          1     1,128.0     1,128.0     1,128    1,128     0.0    pthread_cond_broadcast
```

This will give you another breakdown of all the API calls and other stats that occurred during your program's runtime.

This command does output two files, but the information in said files is not practical for our purposes.